



# arm

## Accelerating oneDNN on AArch64

Jonathan Deakin, Milos Puzovic and team  
5 June 2024

# arm

## Introduction

# Acceleration on CPU: basics

## + Scalar -> NEON/SVE SIMD

- NEON is 128bit, ~4x throughput for f32 (per instruction), ~16x for i8
- SVE has variable vector length, enabling vector length agnostic programming

## + Multi-threading

- Theoretically ~number of core x speedup
- Practically always a bit less, sometimes a lot less

## + Memory access

- Data reuse
- Contiguous access, or packing for specific instructions

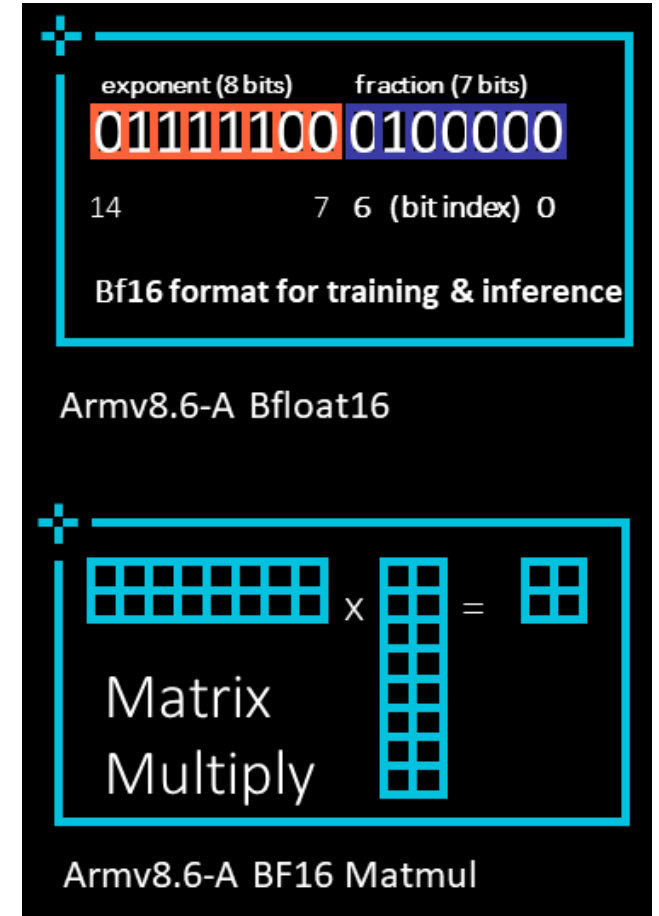
## + Quantization: convert f32 -> i8 => ~4x speedup

## + Fast math: convert f32 -> bf16 => ~2x speedup

## + Function approximation, particularly for exp-based activations

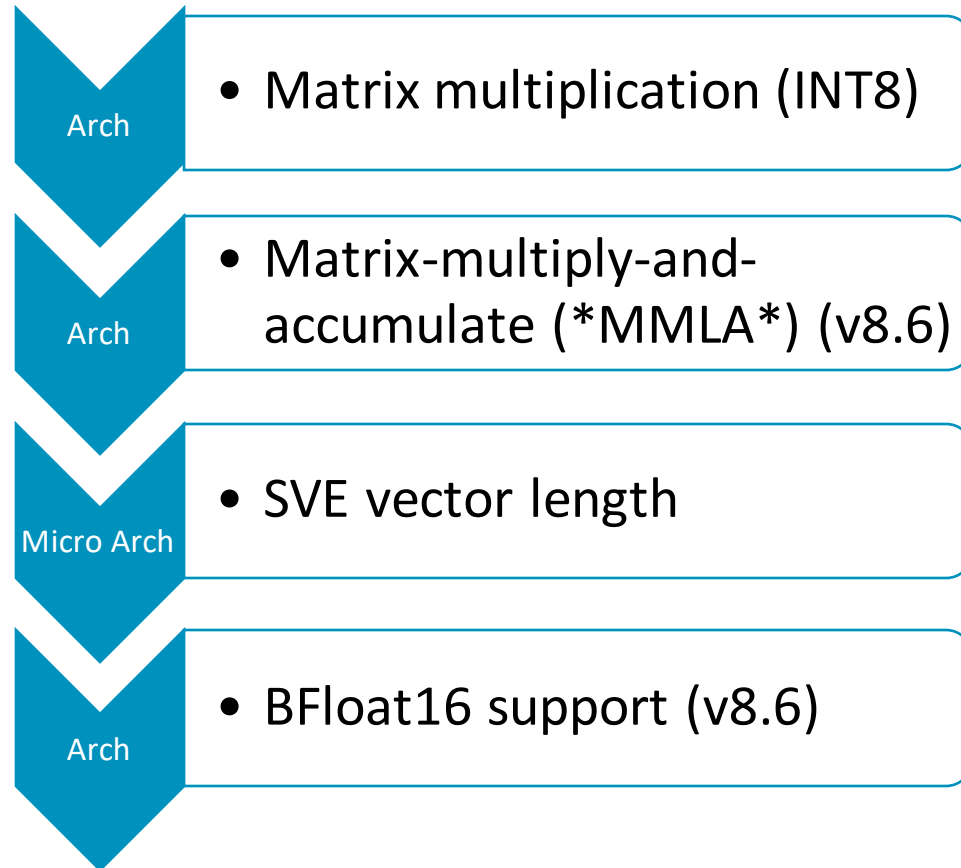
# BFloat16 Extension

- + Armv8.6-A extension, mandatory in Armv9-A
- + Four new instructions
- + Available in SVE, Neon (AArch64 and AArch32)
- + Accepts BF16 inputs, does not generate BF16 results
- + Accumulating an FP32 intermediate result for better accuracy



# ML Specific features in Arm Neoverse cores

## Features enhancing ML performance on Arm Neoverse CPUs



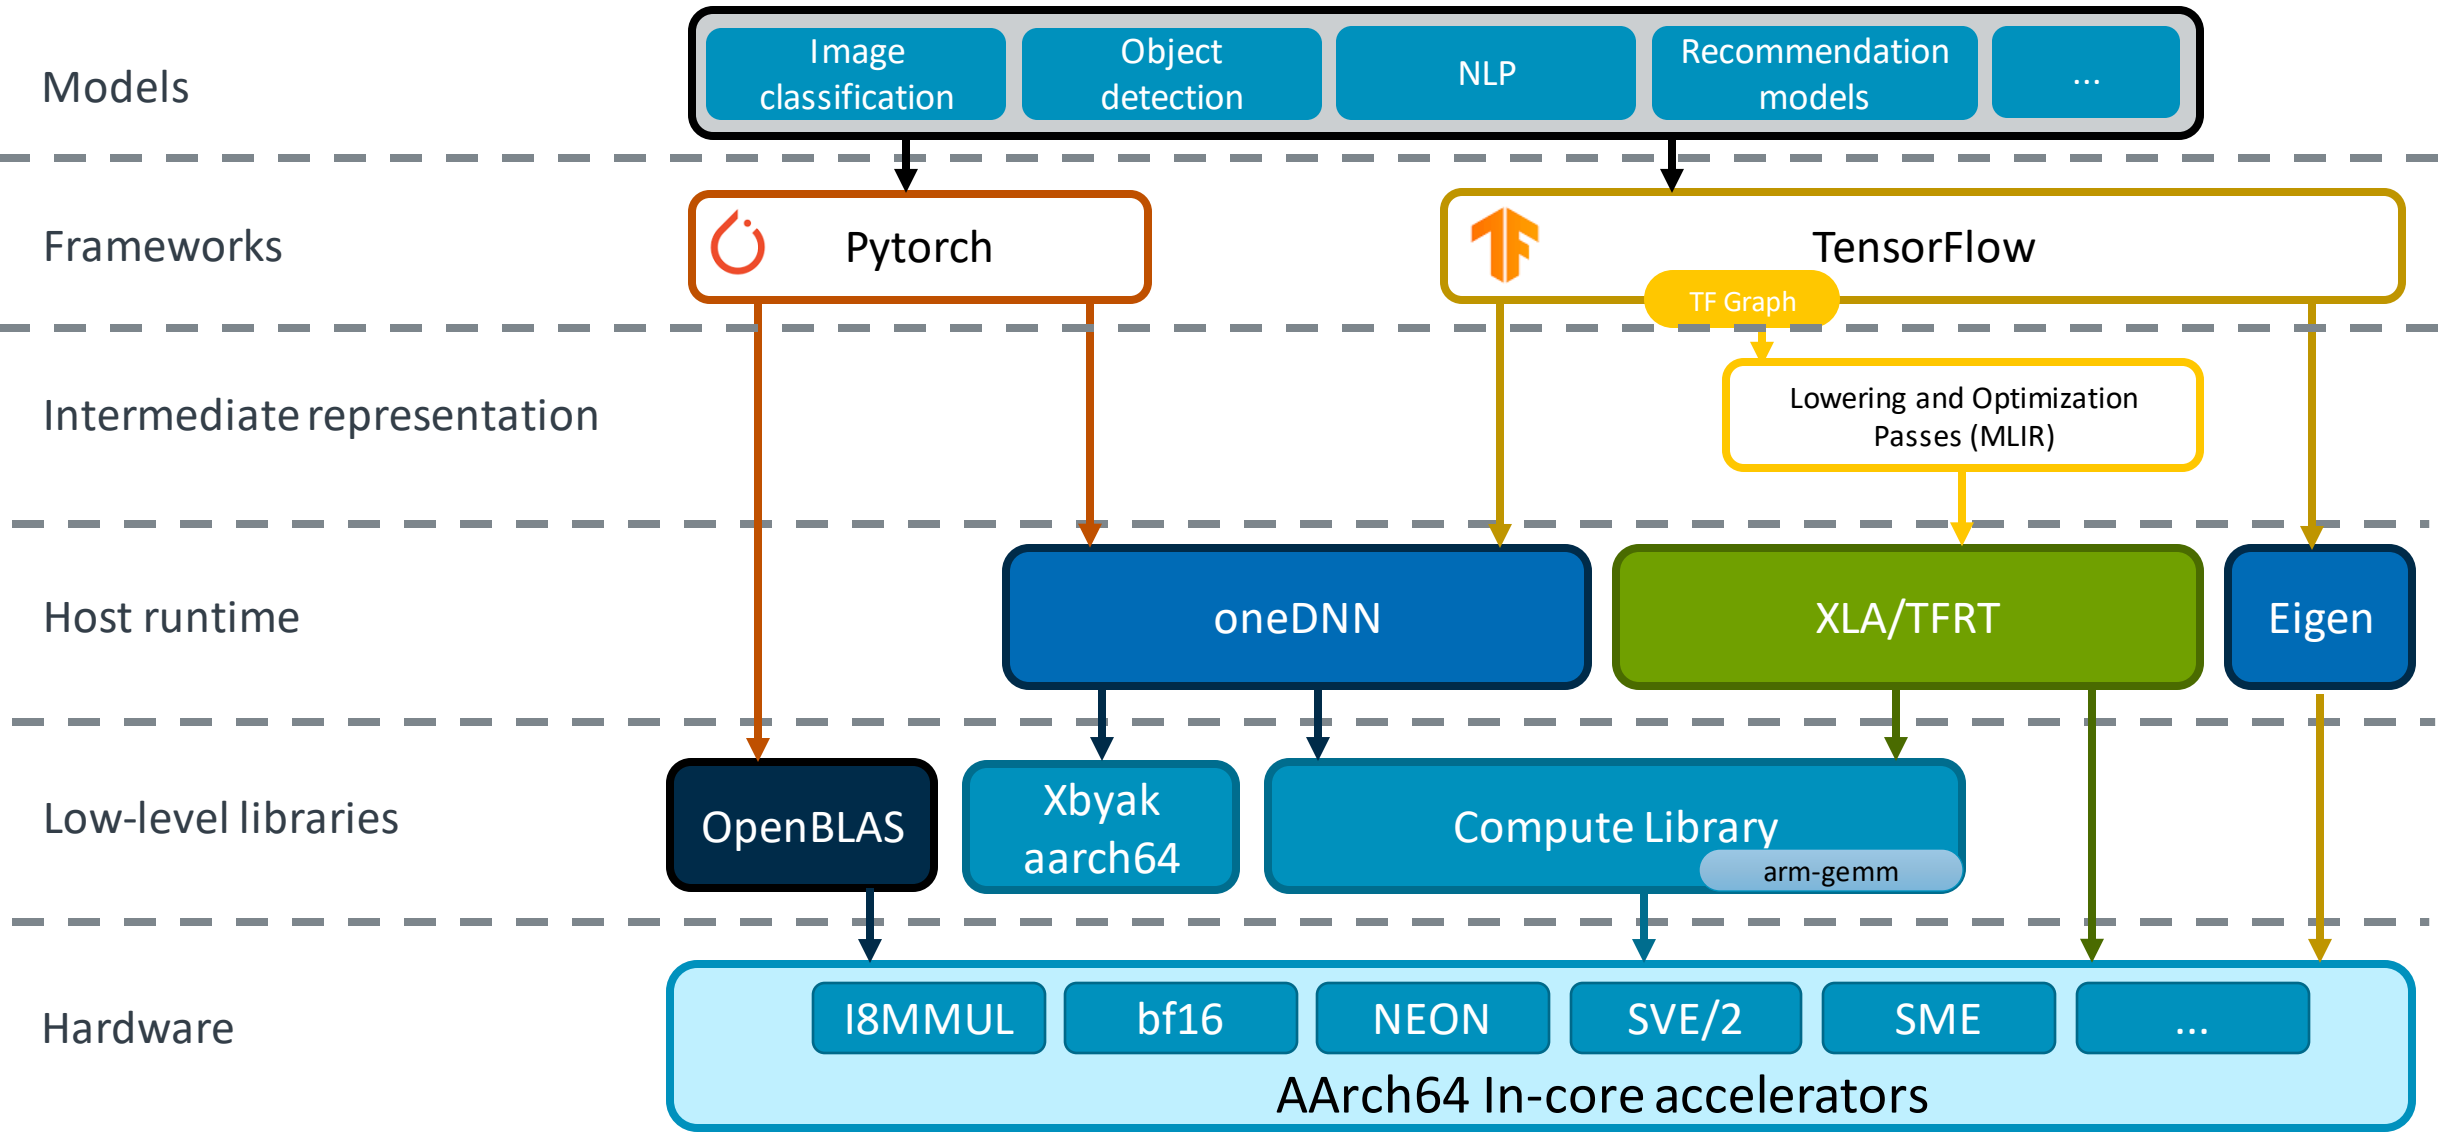
	Neoverse-N1	Neoverse-V1	Neoverse-N2	Neoverse-V2
Vector units	NEON 128-bit x 2	NEON 128-bit x 4 SVE 256-bit x 2	NEON 128-bit x 2 SVE 128-bit x 2	NEON 128-bit x 4 SVE 128-bit x 4
FP32	✓	✓	✓	✓
FP16	✓	✓	✓	✓
BF16	✗	✓	✓	✓
Dot Product	✓	✓	✓	✓
Integer Matrix Multiplier (I8MM)	✗	✓	✓	✓

# MACs/cycle in Arm Neoverse Cores

	Neoverse-N1	Neoverse-V1	Neoverse-N2	Neoverse-V2
<b>FP32 MACs/cycle</b>	8	16	8	16
<b>FP16 MACs/cycle</b>	16	32	16	32
<b>BF16 MACs/cycle</b>	N/A	32 (BFDOT) 64 (BFMMLA)	16 (BFDOT) 32 (BFMMLA)	32 (BFDOT) 64 (BFMMLA)
<b>INT8 MACs/cycle</b>	32 (DOT)	64 (DOT) 128 (I8MM)	32 (DOT) 64 (I8MM)	64 (DOT) 128 (I8MM)

Note: FLOP/cycle (or OP/cycle for int8) = 2 x MACs/cycle x clock speed

# Software stack





# arm

## Previous work



# The Arm Compute Library

Optimized low-level kernels for Arm CPUs

## Key Features:

- Over 100 machine learning functions
- Targeting Arm CPUs and GPUs
- Multiple convolution algorithms (GEMM, Winograd, FFT and Direct)
- Multiple data types: f32, f16, s8, u8, bf16
- Micro-architecture optimization for key ML primitives
- Highly configurable build options enabling lightweight binaries
- Open-source, via permissive MIT license

<https://github.com/ARM-software/ComputeLibrary>

# oneDNN + ACL

- + Accelerated primitives in oneDNN using ACL
  - Matmul + inner product
  - Convolutions (GEMM, indirect, Winograd, depthwise)
  - Eltwise
  - Reorder
  - Softmax
  - Batch/layer norm
  - Binary
- + `acl_post_ops_t` for matmul/ip/conv as a fallback for unfusable post ops
- + “Fixed format” weights allows for reorders to be only done once
  - Using `TF_ONEDNN_ASSUME_FROZEN_WEIGHTS=1`

# xbyak\_aarch64

## + JIT assembler for AArch64 CPUs by C++

“Xbyak\_aarch64 is a C++ library which enables run-time assemble coding with the AArch64 instruction set of Arm(R)v8-A architecture. Xbyak\_aarch64 is based on Xbyak developed for x86\_64 CPUs by MITSUNARI Shigeo.”

## + Use within oneDNN predominantly developed by Fujitsu to accelerate A64FX cores (512bit SVE)

## + Accelerated primitives in oneDNN


- Conv (only 512bit SVE)
- Matmul (using brgemm)
- Reorders
- Softmax
- Pooling
- Eltwise

[https://github.com/fujitsu/xbyak\\_aarch64](https://github.com/fujitsu/xbyak_aarch64)


# Performance improvements



Performance improvements on subset of image classification models  
(normalized to baseline without optimisations)  
**c7g.xlarge, 16 threads, batch 1**



Performance progression on models from MLPerf Suite  
(normalized to baseline without optimisation)  
**c7g.xlarge, 16 threads, batch 8**



# arm

## Ongoing work



# Caching and memory allocation: “stateless”

- + The oneDNN + ACL operators currently
  - Allocate their own memory
  - Store pointers to tensors
- + They have state which is **not captured by the problem descriptor**
- + This is why we must use the resource mapper and a mutex
- + Ongoing work
  - Work is ongoing to expose a “stateless” ACL
  - Move allocations to oneDNN scratchpad
- + This should allow us
  - To make use of oneDNN caching => reduce tech debt
  - Allow concurrent use of a primitive => reduce startup cost
  - Allow sharing of temporary => reduce allocations

# Dynamic quantization

- + Quantize f32 to i8 at runtime based on actual inputs
- + Suitable s8:s8:f32 matmul kernels in ACL 24.04 and oneDNN v3.5
  - >10x faster than gemm:jit
  - ~4x faster than acl F32
- + Pending PRs in iDeep and PyTorch to enable their use
- + Ongoing work to optimize for smaller problems and high thread counts



# Static quantization

- + Quantize f32 to i8 with scales and shifts fixed at init time based on input distributions
- + Much harder than dynamic! User story for static quantization is still unclear in PyTorch
- + Why bother? Allows you to keep inputs i8 between layers (no dequant/requant)
  - Is this a big difference in practice? Difference is not  $O(n^3)$ . Even less difference if accumulator is f32
- + To enable this in PyTorch we need u/s8:s8:u/s8 kernels in oneDNN
- + oneDNN requires bias + post ops to be done in f32
- + We are enabling s8:s8:s8 matmul using separate s8:s8:f32 + post ops + quantize
- + We are working on a s8:s8:s8 convolution using fused requantization in ACL (no posts/bias yet)
- + Still not sure if we can do s32 bias

# Other work

- + JITed reorder for bf16 reorders
- + Pure bf16 matmul and convolution (for torch.autocast)
- + Aarch64 brgemm (Fujitsu)
- + Expanded support in JITed kernels (batch norm, pooling)

# arm

## Future work



# KleidiAI

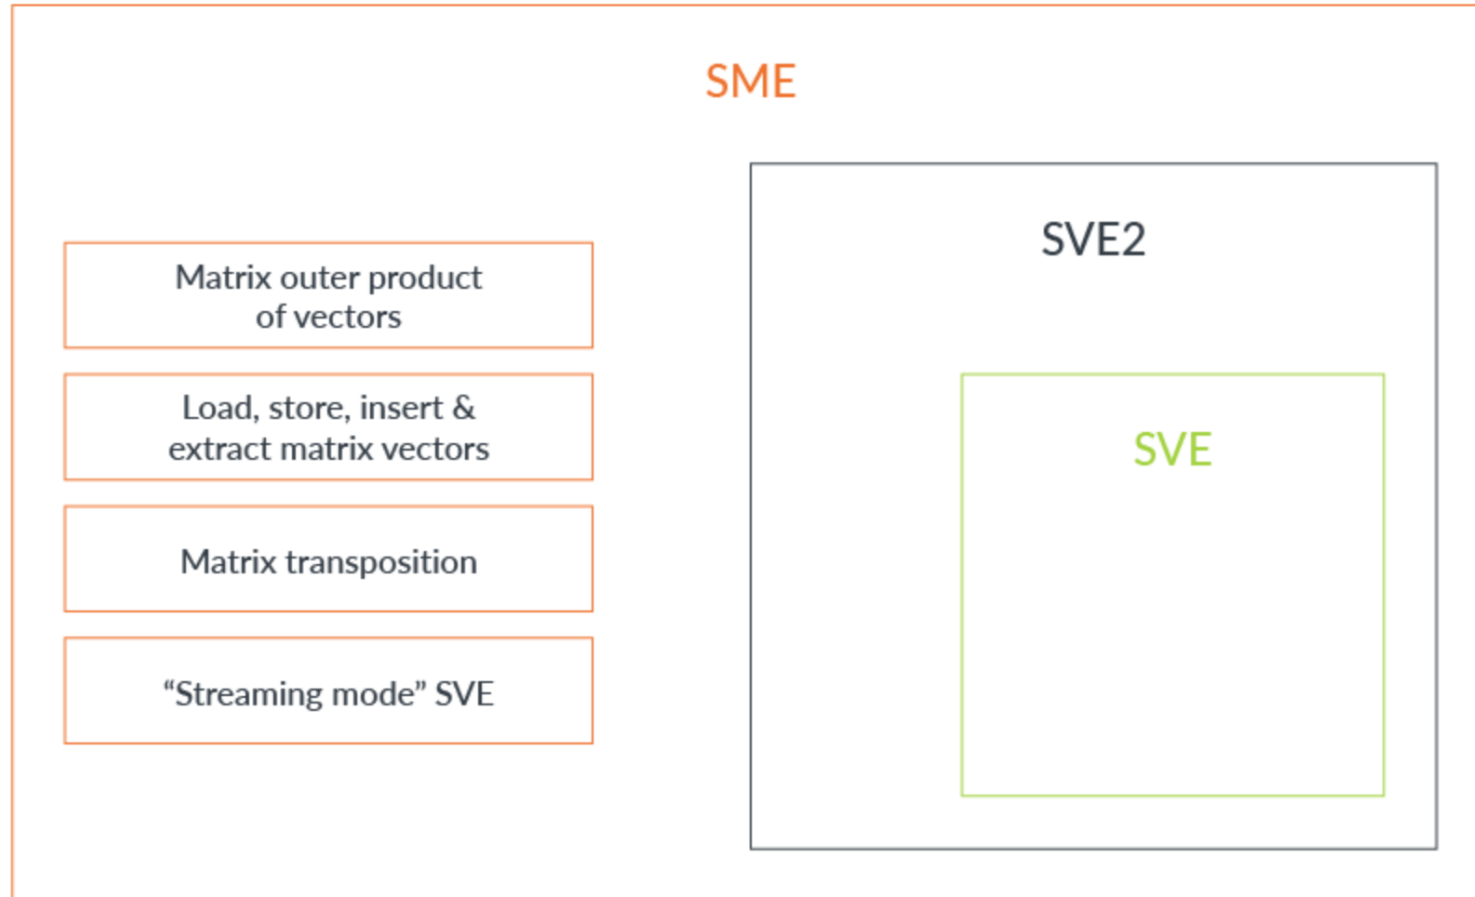
“KleidiAI is an open-source library that provides optimized performance-critical routines, also known as micro-kernels, for artificial intelligence (AI) workloads tailored for Arm<sup>®</sup> CPUs.”

i.e. low level integration

- + ACL was designed as a runtime library (manages threads and memory, has a narrow API)
- + It has allowed us to quickly accelerate a lot of operations
- + ...but as we do more, the oneDNN-ACL API friction is growing
- + Stateless ACL is a step in the right direction
- + But we could go even lower

<https://gitlab.arm.com/kleidi/kleidi.ai>

# Scalable Matrix Extension (SME)



arm

Thank You

Danke

Gracias

Grazie

谢谢

ありがとう

Asante

Merci

감사합니다

धन्यवाद

Kiitos

شكراً

ধন্যবাদ

תודה

ధన్యవాదములు



The Arm trademarks featured in this presentation are registered trademarks or trademarks of Arm Limited (or its subsidiaries) in the US and/or elsewhere. All rights reserved. All other marks featured may be trademarks of their respective owners.

[www.arm.com/company/policies/trademarks](http://www.arm.com/company/policies/trademarks)